

EE58J - Homework 1

Product Image Recognition Challenge 2019

Efehan Atıcı - 2016700186

22 March 2019

1 Size Normalization

In this step I've resized every image in the dataset to 128x128 pixels. This function is called 'resizeBatch' on util.py in my code. We run this resizing algorithm for all folders in the dataset, 'confectionery', 'icecream', 'laundry', 'softdrinks-1' and 'softdrinks-2'.

After we resize all the images to 128x128, we can continue with other steps.

2 Image Description

2.1 Color Histogram

In this step I've created color histograms for every image in the dataset. This function is called 'colorHist' on util.py in my code. It checks for all files ending with '.jpg' extension, loads images and converts them to HSV color space. According to the given window size (1x1 or 2x2 windows for each image), it creates an array that consists of Hue, Saturation and Value histograms. And then we normalize the HSV histograms and append them to a final array which describes the image. Finally, we save this image description as 'crop###.color.npy'. We then use this .npy to save time for calculating color histograms on the next steps.

For creating the histograms, I used bin size as 10 bins for every histogram. Using more than 10 bins makes a big performance hit (finding all descriptors and making classifications for 15 bins with 4x4 windows took more than 30 minutes). Using less than 10 bins would mean there is less information about the image. Choosing 10 bins for histograms seemed the ideal amount, so I used 10 bins when creating every histogram.

Code snippet for this step is as follows:

```
# Only look for files that end with .jpg  
if name.endswith(".jpg"):
```

```

# Load image file
filePath = os.path.join(root, name)
im = cv2.imread(filePath, 1)

# Convert to HSV color space
HSV_im = cv2.cvtColor(im, cv2.COLOR_BGR2HSV)

# Grid stride value for chosen window size
s = int(128 / windowNr)

# Final array for holding all the histograms
full_hsv_hist = []
for i in range(windowNr):
    for j in range(windowNr):

        # Take a window from image
        # given window size
        window=im[i*s:(i+1)*s, j*s:(j+1)*s]

        # Create Hue, Sat, Value Histograms
        hHist=np.histogram(window[:, :, 0])
        sHist=np.histogram(window[:, :, 1])
        vHist=np.histogram(window[:, :, 2])

        # Normalize HSV histograms
        hNorm=np.divide(hHist[0], np.sum(hHist[0]))
        sNorm=np.divide(sHist[0], np.sum(sHist[0]))
        vNorm=np.divide(vHist[0], np.sum(vHist[0]))

        # Concatenate histograms into 1 big array
        hsvHist=np.concatenate((hNorm, sNorm, vNorm))

        # Append them to a final array
        full_hsv_hist.append(hsvHist)
saveName = name.replace(".jpg", "_color.npy")
np.save(os.path.join(root, saveName), full_hsv_hist)

```

After we create color histogram for every image, we move onto the next step.

2.2 Gradient Orientation Histogram

In this step I've created gradient orientation histograms for every image as I did for color histograms. This function is called 'HOGHist' on util.py in my code. It checks for all files ending with '.jpg' extension, loads images as grayscale images. According to the given window size, it creates the image gradients for X and Y

coordinates. We do this by using Sobel operator inside the opencv (cv2.Sobel). Then, after creating gradients for X and Y, we find the gradient orientation simply by taking the arctan of Y and X gradients (np.arctan2). Finally, we save this image description as 'crop###_orient.npy'. Code snippet for this step is as follows:

```

if name.endswith(".jpg"):

    # Load image as grayscale
    filePath = os.path.join(root , name)
    im = cv2.imread(filePath , 0)

    # Find X and Y Gradients
    sobelx = cv2.Sobel(im, cv2.CV_64F, 1, 0, ksize=5)
    sobely = cv2.Sobel(im, cv2.CV_64F, 0, 1, ksize=5)

    # Find orientations in angles , from X and Y grad
    orient = np.arctan2(sobely , sobelx) * 180 / np.pi

    for i in range(windowNr):
        for j in range(windowNr):
            ...

            # Find the histogram for the current window
            orient_hist=np.histogram(window , bins=bin_num)

            # L1 Normalize the histogram
            ...

            # Save the descriptor for future
            saveName = name.replace(".jpg" , "_orient.npy")
            np.save(os.path.join(root , saveName) , full_orient)

```

After we create gradient orientation histogram for every image, we move onto the next step.

3 Nearest Neighbor (NN) Classification

In this step, after creating color and gradient orientation histograms for every image, I've separated images into training and test sets by 80% and 20% respectively.

In total there are 10362 images in the dataset. Among these images, I've randomly selected 2037 images as test set, and the rest 8325 images as training set. I've selected roughly 20% of images from all classes as test set.

3.1 NN Classification using Color Histograms

This function is called ‘nnColor’ in util.py in my code. It takes the all color histograms inside the training dataset, and then appends them to create big array consisting of all the training color histograms. After that, we check for every image in the test set and subtract the color histogram from every element in the training array. We then take the absolute value of this subtracted array of histograms, and find the element with the minimum mean value which tells us the closest neighbour to our guessed image.

Using this technique, we get the following accuracy:

Window Size	Accuracy	Correct	Total
1x1	38.49%	784	2037
2x2	44.58%	908	2037
4x4	47.72%	972	2037

As the window size increases, we see an increase in the accuracy. This happens because there is more information contained in the descriptor. There will be 16 different histograms for 4x4 window size, but there will be only 1 histogram that explains the image in the 1x1 window.

Here is the table of accuracy for every product category when window size is 1x1:

Category	Accuracy	Correct	Total
Confectionery	31.94%	130	407
Icecream	27.38%	115	420
Laundry	45.70%	186	407
Softdrinks-I	23.81%	95	399
Softdrinks-II	58.42%	236	404

On 1x1 window size, using the color histograms descriptors, the most confused class is Softdrinks-I (23.81% accuracy), followed closely by Icecream (27.38%). However the most accurately predicted product category is Softdrinks-II (58.42%), with over 50% accuracy.

Here is the table of accuracy for every product category when window size is 2x2:

Category	Accuracy	Correct	Total
Confectionery	37.59%	153	407
Icecream	31.67%	133	420
Laundry	55.28%	226	407
Softdrinks-I	29.82%	119	399
Softdrinks-II	60.89%	246	404

On 2x2 window size, using the color histograms descriptors, the most confused class is Softdrinks-I (29.82%). The most accurately predicted product category is Softdrinks-II (60.89%).

Here is the table of accuracy for every product category when window size is 4x4:

Category	Accuracy	Correct	Total
Confectionery	40.79%	166	407
Icecream	40.24%	169	420
Laundry	60.93%	248	407
Softdrinks-I	31.58%	126	399
Softdrinks-II	67.33%	272	404

On 4x4 window size, we see accuracy increase in every category. Still, the worst guessed category is Softdrinks-I (31.58%) and the best guessed category is Softdrinks-II (67.33%).




Figure 1: Plot for accuracy of every window size using color histograms

3.2 NN Classification - Gradient Orientation Histograms

This function is called 'nnOrient' in util.py in my code. It works the same way as calculating the classification for Color Histograms.

Using this technique, we get the following accuracy:

Window Size	Accuracy	Correct	Total
1x1	8.15%	166	2037
2x2	26.95%	549	2037
4x4	40.89%	833	2037

As we can see, we get lower accuracy than using the color histograms method. However, as the window size increases, the accuracy gets closer to the color histograms method.

Here is the table of accuracy for every product category when window size is 1x1:

Window Size	Accuracy	Correct	Total
Confectionery	7.62%	31	407
Icecream	8.09%	34	420
Laundry	12.78%	52	407
Softdrinks-I	5.01%	20	399
Softdrinks-II	10.40%	42	404

On 1x1 window size, the most confused class is Softdrinks-I (5.01%). Whereas the most accurately predicted product category is Laundry (12.78%).

Here is the table of accuracy for every product category when window size is 2x2:

Category	Accuracy	Correct	Total
Confectionery	20.39%	83	407
Icecream	19.52%	82	420
Laundry	40.54%	165	407
Softdrinks-I	20.30%	81	399
Softdrinks-II	31.19%	126	404

On 2x2 window size, bad accuracy shifts to the Icecream product category. However, Icecream, Confectionery and Softdrinks-I classes have almost equal accuracies. All these categories perform worse than Laundry and Softdrinks-II.

Here is the table of accuracy for every product category when window size is 4x4:

Category	Accuracy	Correct	Total
Confectionery	30.95%	126	407
Icecream	33.33%	140	420
Laundry	48.16%	196	407
Softdrinks-I	28.82%	115	399
Softdrinks-II	53.96%	218	404

On 4x4 window size, we see accuracy increase in every category. Still, the worst guessed category is Softdrinks-I (28.82%) and the best guessed category is Softdrinks-II (53.96%).

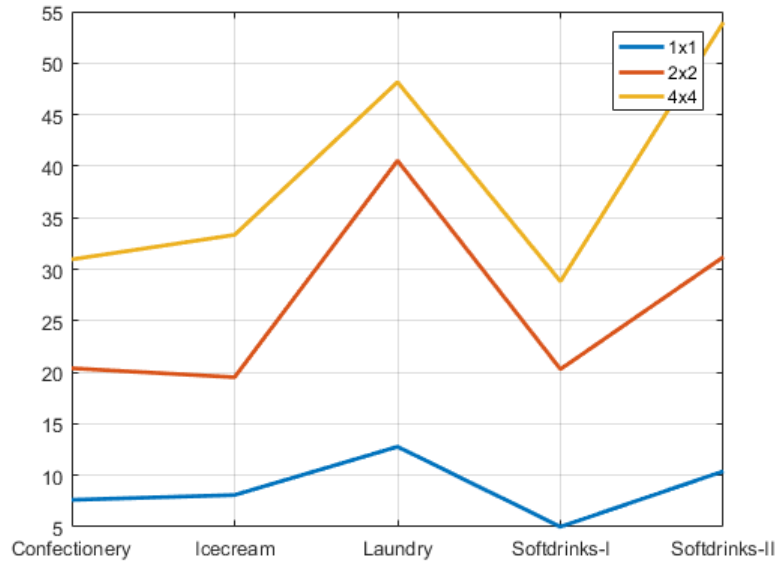


Figure 2: Plot for accuracy of every window size using gradient orientations

3.3 NN Classification - Combination

This function is called ‘nnCombine’ in util.py in my code. In this function, we calculate the difference between color histograms and orientation histograms separately, and then we sum them up to find the closest neighbour. Using this way, we get a higher accuracy than both of the other techniques:

Window Size	Accuracy	Correct	Total
1x1	42.02%	856	2037
2x2	51.55%	1050	2037
4x4	56.90%	1159	2037

Since we combined both techniques, we get a higher accuracy for detection overall. Our best classification accuracy is 56.90% with 1159 correct guesses out of 2037 total images when we have 4x4 windows on the image.

Here is the table of accuracy for every product category when window size is 1x1:

Category	Accuracy	Correct	Total
Confectionery	37.84%	154	407
Icecream	34.05%	143	420
Laundry	48.89%	199	407
Softdrinks-I	25.06%	100	399
Softdrinks-II	64.36%	260	404

On 1x1 window size, using the combination of descriptors, the most confused class is Softdrinks-I with 25.06% accuracy. Whereas the most accurately predicted product category is Softdrinks-II with 64.36% accuracy.

Here is the table of accuracy for every product category when window size is 2x2:

Category	Accuracy	Correct	Total
Confectionery	45.70%	186	407
Icecream	40.00%	168	420
Laundry	59.95%	244	407
Softdrinks-I	39.60%	158	399
Softdrinks-II	71.78%	290	404

On 2x2 window size, both Icecream and Softdrinks-I categories share bad accuracy. However, Softdrinks-II category is the best performing category, with over 70% accuracy!

Here is the table of accuracy for every product category when window size is 4x4:

Category	Accuracy	Correct	Total
Confectionery	45.20%	184	407
Icecream	46.19%	194	420
Laundry	60.44%	246	407
Softdrinks-I	42.35%	169	399
Softdrinks-II	78.21%	316	404

On 4x4 window size, we see a little accuracy drop in confectionery, and we don't see any big improvements on other categories. Still, the worst guessed category is Softdrinks-I and the best guessed category is Softdrinks-II.

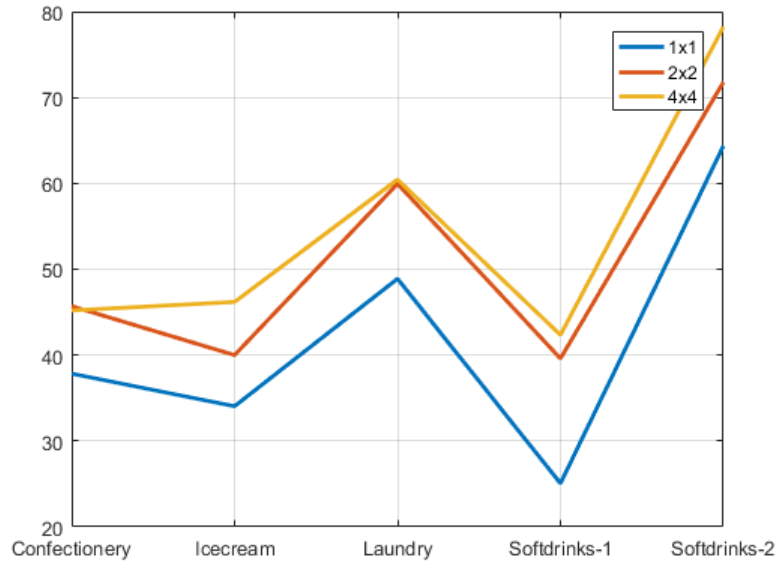


Figure 3: Plot for accuracy of every window size using both descriptors